

Introduction to Aerial Robotics

Lecture 9

Shaojie Shen

Assistant Professor

Dept. of ECE, HKUST



10 November 2015

Outline

- Extended Kalman Filter
- Particle Filter

Extended Kalman Filter

Bayes' Filter

- **Prior:** $p(x_0)$ ← State
- **Process model:** $f(x_t | x_{t-1}, u_t)$ ← Control input
- **Measurement model:** $g(z_t | x_t)$ ← Measurement
- **Prediction step:**
- $p(x_t | z_{1:t-1}, u_{1:t}) = \int f(x_t | x_{t-1}, u_t) p(x_{t-1} | z_{1:t-1}, u_{1:t-1}) dx_{t-1}$
- **Update step:**
- $$p(x_t | z_{1:t}, u_{1:t}) = \frac{g(z_t | x_t) p(x_t | z_{1:t-1}, u_{1:t})}{\int g(z_t | x'_t) p(x'_t | z_{1:t-1}, u_{1:t}) dx'_t}$$

Assumptions

- The prior state of the robot is represented by a Gaussian distribution
 - $p(x_0) \sim N(\mu_0, \Sigma_0)$
- The process model $f(x_t | x_{t-1}, u_t)$ is linear with additive Gaussian white noise
 - $x_t = A_t x_{t-1} + B_t u_t + n_t$
 - $n_t \sim N(0, Q_t)$
- The measurement model $g(z_t | x_t)$ is linear with additive Gaussian white noise
 - $z_t = C_t x_t + v_t$
 - $v_t \sim N(0, R_t)$

Kalman Filter

- Prior:
 - $p(x_0) \sim N(\mu_0, \Sigma_0)$
- Transition model:
 - $x_t = A_t x_{t-1} + B_t u_t + n_t$
 - $n_t \sim N(0, Q_t)$
- Measurement model:
 - $z_t = C_t x_t + v_t$
 - $v_t \sim N(0, R_t)$
- Prior:
 - μ_{t-1}, Σ_{t-1}
- Prediction:
 - $\bar{\mu}_t = A_t \mu_{t-1} + B_t u_t$
 - $\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + Q_t$
- Update:
 - $\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$
 - $\Sigma_t = \bar{\Sigma}_t - K_t C_t \bar{\Sigma}_t$
 - $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + R_t)^{-1}$

Example Problem

$$x_t = x_{t-1} + u_t + n_t$$


$$Q_t = 0.5$$

$$A_t = B_t = 1$$

$$z_t = x_t + v_t$$

$$R_t = 1.0$$

$$C_t = 1$$




Continuous Dynamics

- Can convert continuous time systems
- $\dot{x} = f(x, u, n) = A x + B u + U n$
- Into discrete time systems using one-step Euler integration
- $x_t = F x_{t-1} + G u_t + V n_t$
- $F = (I + \delta t A), G = \delta t B, V = \delta t U$
- This will introduce some error, but the observations can help correct it
- Prediction:
 - $\bar{\mu}_t = F \mu_{t-1} + G u_t$
 - $\bar{\Sigma}_t = F \Sigma_{t-1} F^T + V Q V^T$

Kalman Filter Discussion

- **Advantages:**
 - Simple
 - Purely matrix operations
 - Computationally efficient, even for high dimensional systems
- **Disadvantages:**
 - Assumes everything is linear and Gaussian
 - Unimodal distribution
 - Cannot handle multiple hypotheses



Extended Kalman Filter (EKF)

Assumptions


- The prior state of the robot is represented by a Gaussian distribution
 - $p(x_0) \sim N(\mu_0, \Sigma_0)$
- The continuous time process model is:
 - $\dot{x} = f(x, u, n)$
 - $n_t \sim N(0, Q_t)$ is Gaussian white noise
- The observation model is:
 - $z = h(x, v)$
 - $v_t \sim N(0, R_t)$ is Gaussian white noise

Prediction

- Process model is nonlinear
- Need to convert the continuous dynamics to a discrete time system
- Look over a finite time interval $\tau = [t', t)$, where $t - t' = \delta t$
 - $t' \rightarrow t - 1$, \bar{t} is an infinitesimal step before t
- Options:
 - Integrate the process model over the time horizon τ
 - $x_{\bar{t}} = \Phi(\bar{t}; x_{t-1}, u, n)$
 - Difficult to do in general
 - Use numerical integration
 - One-step Euler integration

Prediction – Linearization

- Linearize the dynamics about $x = \mu_{t-1}$, $u = u_t$, $n = 0$
 - $$\dot{x} \approx f(\mu_{t-1}, u_t, 0) + \left. \frac{\partial f}{\partial x} \right|_{\mu_{t-1}, u_t, 0} (x - \mu_{t-1}) + \left. \frac{\partial f}{\partial u} \right|_{\mu_{t-1}, u_t, 0} (u - u_t) + \left. \frac{\partial f}{\partial n} \right|_{\mu_{t-1}, u_t, 0} (n - 0)$$
- Let:
 - $$A_t = \left. \frac{\partial f}{\partial x} \right|_{\mu_{t-1}, u_t, 0}$$
 - $$B_t = \left. \frac{\partial f}{\partial u} \right|_{\mu_{t-1}, u_t, 0}$$
 - $$U_t = \left. \frac{\partial f}{\partial n} \right|_{\mu_{t-1}, u_t, 0}$$
- Linear dynamics:
 - $$\dot{x} \approx f(\mu_{t-1}, u_t, 0) + A_t(x - \mu_{t-1}) + B_t(u - u_t) + U_t(n - 0)$$



Prediction – Discrete Time

- One-step Euler integration

- $x_{\bar{t}} \approx x_{t-1} + f(x_{t-1}, u_t, n_t) \delta t$

- $\approx x_{t-1} + \delta t f(\mu_{t-1}, u_t, 0) + \delta t A_t (x_{t-1} - \mu_{t-1}) + \delta t B_t(u_t - u_t)$
 $+ \delta t U_t(n_t - 0)$

- $\approx x_{t-1} + \delta t f(\mu_{t-1}, u_t, 0) + \delta t A_t (x_{t-1} - \mu_{t-1}) + \delta t U_t(n_t - 0)$

- $\approx (I + \delta t A_t) x_{t-1} + \delta t U_t n_t + \delta t(f(\mu_{t-1}, u_t, 0) - A_t \mu_{t-1})$

- $\approx F_t x_{t-1} + V_t n_t + \delta t(f(\mu_{t-1}, u_t, 0) - A_t \mu_{t-1})$

- Prediction:

- $\bar{\mu}_t = \mu_{t-1} + \delta t f(\mu_{t-1}, u_t, 0)$

- $\bar{\Sigma}_t = F_t \Sigma_{t-1} F_t^T + V_t Q_t V_t^T$

Update – Linearization

- Linearize the observation model about $x = \bar{\mu}_t$, $v = 0$
 - $g(x, v) \approx g(\bar{\mu}_t, 0) + \left. \frac{\partial g}{\partial x} \right|_{\bar{\mu}_t, 0} (x - \bar{\mu}_t) + \left. \frac{\partial g}{\partial v} \right|_{\bar{\mu}_t, 0} (v - 0)$
- Let:
 - $C_t = \left. \frac{\partial g}{\partial x} \right|_{\bar{\mu}_t, 0}$
 - $W_t = \left. \frac{\partial g}{\partial v} \right|_{\bar{\mu}_t, 0}$
- Linear observation model:
 - $z_t = g(x_t, v_t) \approx g(\bar{\mu}_t, 0) + C_t (x_t - \bar{\mu}_t) + W_t v_t$

Update

- Follow the same derivation as the Kalman Filter
- $$\begin{bmatrix} x_t \\ z_t \end{bmatrix} = \begin{bmatrix} I & 0 \\ C_t & W_t \end{bmatrix} \begin{bmatrix} \bar{x}_t \\ v_t \end{bmatrix} + \begin{bmatrix} 0 \\ g(\bar{\mu}_t, 0) - C_t \bar{\mu}_t \end{bmatrix}$$
- Mean:
 - $E[X_t] = E[\bar{X}_t] = \bar{\mu}_t$
 - $E[Z_t] = E[C_t \bar{X}_t + W_t V_t + g(\bar{\mu}_t, 0) - C_t \bar{\mu}_t]$
 - $= C_t \bar{\mu}_t + g(\bar{\mu}_t, 0) - C_t \bar{\mu}_t$
 - $= g(\bar{\mu}_t, 0)$

Update

- Follow the same derivation as the Kalman Filter

- $$\begin{bmatrix} x_t \\ z_t \end{bmatrix} = \begin{bmatrix} I & 0 \\ C_t & W_t \end{bmatrix} \begin{bmatrix} x_{\bar{t}} \\ v_t \end{bmatrix} + \begin{bmatrix} 0 \\ g(\bar{\mu}_t, 0) - C_t \bar{\mu}_t \end{bmatrix}$$

- Covariance:

$$\begin{aligned} - \Sigma &= \begin{bmatrix} I & 0 \\ C_t & W_t \end{bmatrix} \begin{bmatrix} \bar{\Sigma}_t & 0 \\ 0 & R_t \end{bmatrix} \begin{bmatrix} I & C_t^T \\ 0 & W_t^T \end{bmatrix} \\ - &= \begin{bmatrix} \bar{\Sigma}_t & \bar{\Sigma}_t C_t^T \\ C_t \bar{\Sigma}_t & C_t \bar{\Sigma}_t C_t^T + W_t R_t W_t^T \end{bmatrix} \end{aligned}$$

Update

- Recall that for a multivariate Gaussian $Y = \begin{bmatrix} X \\ Z \end{bmatrix}$ with mean $\mu = \begin{bmatrix} \mu_X \\ \mu_Z \end{bmatrix}$ and covariance $\Sigma = \begin{bmatrix} \Sigma_{XX} & \Sigma_{XZ} \\ \Sigma_{ZX} & \Sigma_{ZZ} \end{bmatrix}$
- The conditional density $f_{X|Z}(x | Z = z)$ is Gaussian with
 - $\mu_{X|Z} = \mu_X + \Sigma_{XZ} \Sigma_{ZZ}^{-1} (z - \mu_Z)$
 - $\Sigma_{X|Z} = \Sigma_{XX} - \Sigma_{XZ} \Sigma_{ZZ}^{-1} \Sigma_{ZX}$
- Result:
 - $\mu_t = \bar{\mu}_t + K_t (z_t - g(\bar{\mu}_t, 0))$
 - $\Sigma_t = \bar{\Sigma}_t - K_t C_t^T \bar{\Sigma}_t$
 - $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + W_t R W_t^T)^{-1}$

Extended Kalman Filter

- Prediction step:

- $\bar{\mu}_t = \mu_{t-1} + \delta t f(\mu_{t-1}, u_t, 0)$
- $\bar{\Sigma}_t = F_t \Sigma_{t-1} F_t^T + V_t Q_t V_t^T$
- $\dot{x} = f(x, u, n)$
- $n_t \sim N(0, Q_t)$ } Assumptions
- $A_t = \left. \frac{\partial f}{\partial x} \right|_{\mu_{t-1}, u_t, 0}$
- $U_t = \left. \frac{\partial f}{\partial n} \right|_{\mu_{t-1}, u_t, 0}$ } Linearization
- $F_t = I + \delta t A_t$
- $V_t = \delta t U_t$ } Discretization

- Update step:

- $\mu_t = \bar{\mu}_t + K_t (z_t - g(\bar{\mu}_t, 0))$
- $\Sigma_t = \bar{\Sigma}_t - K_t C_t \bar{\Sigma}_t$
- $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + W_t R W_t^T)^{-1}$
- $z_t = g(x_t, v_t)$
- $v_t \sim N(0, R_t)$ } Assumptions
- $C_t = \left. \frac{\partial g}{\partial x} \right|_{\bar{\mu}_t, 0}$
- $W_t = \left. \frac{\partial g}{\partial v} \right|_{\bar{\mu}_t, 0}$ } Linearization

Example Problem

Quadrotor with a Good Velocity Sensor



State

- Can accurately estimate the commanded linear velocity using the motion tracking system and the angular velocity using a gyroscope

- $\mathbf{x} = \begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \mathbf{x}_3 \end{bmatrix} = \begin{bmatrix} \mathbf{p} \\ \mathbf{q} \\ \mathbf{b}_g \end{bmatrix} = \begin{bmatrix} \text{position} \\ \text{orientation} \\ \text{gyroscope bias} \end{bmatrix} \in \mathbf{R}^9$

- Use Z-X-Y Euler angle parameterization of $SO(3)$ for orientation

- $\mathbf{q} = [\phi, \theta, \psi]^T = [\text{roll}, \text{pitch}, \text{yaw}]^T$

- $R = \begin{bmatrix} c\psi c\theta - s\phi s\psi s\theta & -c\phi s\psi & c\psi s\theta + c\theta s\phi s\psi \\ c\theta s\psi + c\psi s\phi s\theta & c\phi c\psi & s\psi s\theta - c\psi c\theta s\phi \\ -c\phi s\theta & s\phi & c\phi c\theta \end{bmatrix}$

Process Model

- **Assumption:** the motion tracking system gives a noisy estimate of the linear velocity
 - $\mathbf{v}_m = \dot{\mathbf{p}} + \mathbf{n}_v$
- **Assumption:** the gyroscope gives a noisy estimate of the angular velocity
 - $\omega_m = \omega + \mathbf{b}_g + \mathbf{n}_g$
- **Assumption:** the drift in the gyroscope bias is described by a Gaussian, white noise process
 - $\dot{\mathbf{b}}_g = \mathbf{n}_{bg}$
 - $\mathbf{n}_{bg} \sim N(0, Q_g)$

Process Model

- ω_m is in the body frame, \mathbf{q} is in the world frame
- **Recall:** the angular velocity in the body frame is given by

- $$\omega = \begin{bmatrix} p \\ q \\ r \end{bmatrix} = \begin{bmatrix} c\theta & 0 & -c\phi s\theta \\ 0 & 1 & s\phi \\ s\theta & 0 & c\phi c\theta \end{bmatrix} \begin{bmatrix} \dot{\phi} \\ \dot{\theta} \\ \dot{\psi} \end{bmatrix} = G(\mathbf{q})\dot{\mathbf{q}}$$

- Process model:

- $$\dot{\mathbf{x}} = \begin{bmatrix} \mathbf{v}_m - \mathbf{n}_v \\ G(\mathbf{x}_2)^{-1}(\omega_m - \mathbf{x}_3 - \mathbf{n}_g) \\ \mathbf{n}_{bg} \end{bmatrix}$$

Observation Model

- Use a camera to measure the pose of the robot
- Use theory of projective geometry
 - Covered earlier in the semester
- Can estimate the position and orientation of the robot using a minimum of 4 features on the ground plane, e.g., AR Tags
 - Can recover \mathbf{q} from the rotation matrix R
- $\mathbf{z} = \begin{bmatrix} \mathbf{p} \\ \mathbf{q} \end{bmatrix} + \mathbf{v}$
$$= \begin{bmatrix} I & 0 & 0 \\ 0 & I & 0 \end{bmatrix} \begin{bmatrix} \mathbf{p} \\ \mathbf{q} \\ \mathbf{b}_g \end{bmatrix} + \mathbf{v}$$
$$= C \mathbf{x} + \mathbf{v}$$

Example Problem

Quadrotor with a Good Acceleration Sensor



State

- Can accurately estimate the commanded linear acceleration and angular velocity using the IMU

- $$\mathbf{x} = \begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \mathbf{x}_3 \\ \mathbf{x}_4 \\ \mathbf{x}_5 \end{bmatrix} = \begin{bmatrix} \mathbf{p} \\ \mathbf{q} \\ \dot{\mathbf{p}} \\ \mathbf{b}_g \\ \mathbf{b}_a \end{bmatrix} = \begin{bmatrix} \text{position} \\ \text{orientation} \\ \text{linear velocity} \\ \text{gyroscope bias} \\ \text{accelerometer bias} \end{bmatrix} \in \mathbf{R}^{15}$$

- Use Z-X-Y Euler angle parameterization of $SO(3)$ for orientation
 - $\mathbf{q} = [\phi, \theta, \psi]^T = [\text{roll}, \text{pitch}, \text{yaw}]^T$

- $$R = \begin{bmatrix} c\psi c\theta - s\phi s\psi s\theta & -c\phi s\psi & c\psi s\theta + c\theta s\phi s\psi \\ c\theta s\psi + c\psi s\phi s\theta & c\phi c\psi & s\psi s\theta - c\psi c\theta s\phi \\ -c\phi s\theta & s\phi & c\phi c\theta \end{bmatrix}$$

Process Model – Gyroscope

- **Assumption:** the gyroscope gives a noisy estimate of the angular velocity
 - $\omega_m = \omega + \mathbf{b}_g + \mathbf{n}_g$
- **Assumption:** the drift in the gyroscope bias is described by a Gaussian, white noise process
 - $\dot{\mathbf{b}}_g = \mathbf{n}_{bg}$
 - $\mathbf{n}_{bg} \sim N(0, Q_g)$

Process Model – Accelerometer

- **Assumption:** the accelerometer gives a noisy estimate of the linear acceleration
 - $\mathbf{a}_m = R(\mathbf{q})^T (\ddot{\mathbf{p}} - \mathbf{g}) + \mathbf{b}_a + \mathbf{n}_a$
- **Assumption:** the drift in the accelerometer bias is described by a Gaussian, white noise process
 - $\dot{\mathbf{b}}_a = \mathbf{n}_{ba}$
 - $\mathbf{n}_{ba} \sim N(0, Q_a)$

Process Model

- Process model:

$$\dot{\mathbf{x}} = \begin{bmatrix} \mathbf{x}_3 \\ G(\mathbf{x}_2)^{-1}(\omega_m - \mathbf{x}_4 - \mathbf{n}_g) \\ \mathbf{g} + R(\mathbf{x}_2)(\mathbf{a}_m - \mathbf{x}_5 - \mathbf{n}_a) \\ \mathbf{n}_{bg} \\ \mathbf{n}_{ba} \end{bmatrix}$$

Observation Model


- Use a camera to measure:
 - The pose of the robot (using AR Tags)
 - The linear velocity (using optical flow)

- $$\mathbf{z} = \begin{bmatrix} \mathbf{p} \\ \mathbf{q} \\ \dot{\mathbf{p}} \end{bmatrix} + \mathbf{v}$$
$$= \begin{bmatrix} I & 0 & 0 & 0 & 0 \\ 0 & I & 0 & 0 & 0 \\ 0 & 0 & I & 0 & 0 \end{bmatrix} \begin{bmatrix} \mathbf{p} \\ \mathbf{q} \\ \dot{\mathbf{p}} \\ \mathbf{b}_g \\ \mathbf{b}_a \end{bmatrix} + \mathbf{v}$$
$$= \mathbf{C} \mathbf{x} + \mathbf{v}$$

Observation Model

- Use a camera to measure:
 - The pose of the robot (using AR Tags)

$$\begin{aligned} \bullet \quad \mathbf{z} &= \begin{bmatrix} \mathbf{p} \\ \mathbf{q} \\ \dot{\mathbf{p}} \end{bmatrix} + \mathbf{v} \\ &= \begin{bmatrix} I & 0 & 0 & 0 & 0 \\ 0 & I & 0 & 0 & 0 \\ 0 & 0 & I & 0 & 0 \end{bmatrix} \begin{bmatrix} \mathbf{p} \\ \mathbf{q} \\ \dot{\mathbf{p}} \\ \mathbf{b}_g \\ \mathbf{b}_a \end{bmatrix} + \mathbf{v} \\ &= \mathbf{C} \mathbf{x} + \mathbf{v} \end{aligned}$$



Recap

Bayes' Filter

- **Prior:** $p(x_0)$ ← State
- **Process model:** $f(x_t | x_{t-1}, u_t)$ ← Control input
- **Measurement model:** $g(z_t | x_t)$ ← Measurement
- **Prediction step:**
- $p(x_t | z_{1:t-1}, u_{1:t}) = \int f(x_t | x_{t-1}, u_t) p(x_{t-1} | z_{1:t-1}, u_{1:t-1}) dx_{t-1}$
- **Update step:**
- $p(x_t | z_{1:t}, u_{1:t}) = \frac{g(z_t | x_t) p(x_t | z_{1:t-1}, u_{1:t})}{\int g(z_t | x'_t) p(x'_t | z_{1:t-1}, u_{1:t}) dx'_t}$

Assumptions

- The prior state of the robot is represented by a Gaussian distribution
 - $p(x_0) \sim N(\mu_0, \Sigma_0)$
- The continuous time process model is:
 - $\dot{x} = f(x, u, n)$
 - $n_t \sim N(0, Q_t)$ is Gaussian white noise
- The observation model is:
 - $z = h(x, v)$
 - $v_t \sim N(0, R_t)$ is Gaussian white noise

Extended Kalman Filter

- Prediction step:


- $\bar{\mu}_t = \mu_{t-1} + \delta t f(\mu_{t-1}, u_t, 0)$
- $\bar{\Sigma}_t = F_t \Sigma_{t-1} F_t^T + V_t Q_t V_t^T$
- $\dot{x} = f(x, u, n)$
- $n_t \sim N(0, Q_t)$ } Assumptions
- $A_t = \left. \frac{\partial f}{\partial x} \right|_{\mu_{t-1}, u_t, 0}$
- $U_t = \left. \frac{\partial f}{\partial n} \right|_{\mu_{t-1}, u_t, 0}$ } Linearization
- $F_t = I + \delta t A_t$
- $V_t = \delta t U_t$ } Discretization

- Update step:

- $\mu_t = \bar{\mu}_t + K_t (z_t - g(\bar{\mu}_t, 0))$
- $\Sigma_t = \bar{\Sigma}_t - K_t C_t \bar{\Sigma}_t$
- $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + W_t R W_t^T)^{-1}$
- $z_t = g(x_t, v_t)$
- $v_t \sim N(0, R_t)$ } Assumptions
- $C_t = \left. \frac{\partial g}{\partial x} \right|_{\bar{\mu}_t, 0}$
- $W_t = \left. \frac{\partial g}{\partial v} \right|_{\bar{\mu}_t, 0}$ } Linearization

EKF Discussion

- **Advantages:**
 - Simple
 - Computationally efficient, even for high dimensional systems
 - Works with generic process and observation models
- **Disadvantages:**
 - Must calculate the Jacobian of the process and observation models
 - Unimodal distribution
 - Cannot handle multiple hypotheses



Other Applications

- Can use Bayesian filtering methods in many domains:
 - Pose estimation
 - Parameter estimation
 - Map building
 - Simultaneous localization and mapping (SLAM)
 - Feature tracking
 - Target tracking


Particle Filter

Probabilistic Robotics

Bayes Filter Implementations

Particle filters


Sample-based Localization (sonar)



Particle Filters

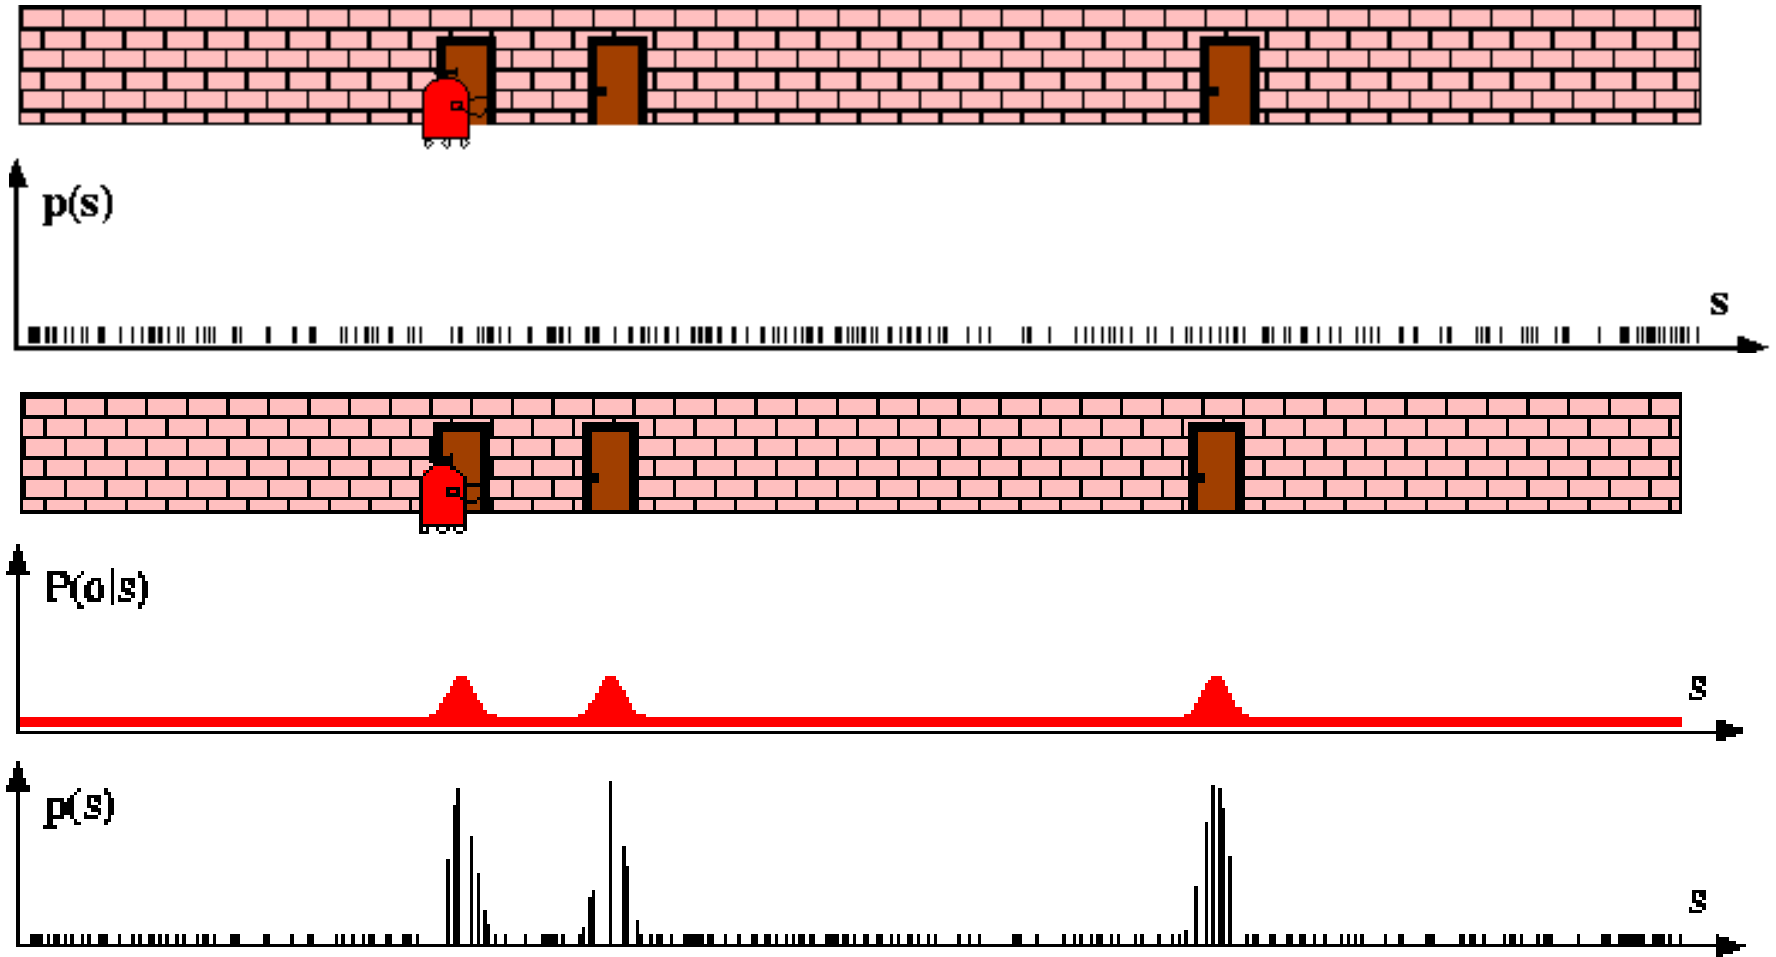
- Represent belief by random **samples**
- Estimation of **non-Gaussian, nonlinear** processes
- Monte Carlo filter, Survival of the fittest, Condensation, Bootstrap filter, Particle filter
- Filtering: [Rubin, 88], [Gordon et al., 93], [Kitagawa 96]
- Computer vision: [Isard and Blake 96, 98]
- Dynamic Bayesian Networks: [Kanazawa et al., 95]

Particle Filters





Sensor Information: Importance Sampling

$$\begin{aligned} Bel(x) &\leftarrow \alpha p(z | x) Bel^-(x) \\ w &\leftarrow \frac{\alpha p(z | x) Bel^-(x)}{Bel^-(x)} = \alpha p(z | x) \end{aligned}$$







Robot Motion

$$Bel^-(x) \leftarrow \int p(x|u, x') Bel(x') dx'$$





Sensor Information: Importance Sampling

$$\begin{aligned} Bel(x) &\leftarrow \alpha p(z|x) Bel^-(x) \\ w &\leftarrow \frac{\alpha p(z|x) Bel^-(x)}{Bel^-(x)} = \alpha p(z|x) \end{aligned}$$



Robot Motion

$$Bel^-(x) \leftarrow \int p(x|u, x') Bel(x') \, dx'$$



Particle Filter Algorithm

1. Algorithm **particle_filter**($S_{t-1}, u_{t-1} z_t$):
2. $S_t = \emptyset, \quad \eta = 0$
3. **For** $i = 1 \dots n$ *Generate new samples*
4. Sample index $j(i)$ from the discrete distribution given by w_{t-1}
5. Sample x_t^i from $p(x_t | x_{t-1}, u_{t-1})$ using $x_{t-1}^{j(i)}$ and u_{t-1}
6. $w_t^i = p(z_t | x_t^i)$ *Compute importance weight*
7. $\eta = \eta + w_t^i$ *Update normalization factor*
8. $S_t = S_t \cup \{ \langle x_t^i, w_t^i \rangle \}$ *Insert*
9. **For** $i = 1 \dots n$
10. $w_t^i = w_t^i / \eta$ *Normalize weights*

Particle Filter Algorithm

$$Bel(x_t) = \eta p(z_t | x_t) \int p(x_t | x_{t-1}, u_{t-1}) Bel(x_{t-1}) dx_{t-1}$$

draw x_{t-1}^i from $Bel(x_{t-1})$

draw x_t^i from $p(x_t | x_{t-1}^i, u_{t-1})$


Importance factor for x_t^i :

$$\begin{aligned} w_t^i &= \frac{\text{target distribution}}{\text{proposal distribution}} \\ &= \frac{\eta p(z_t | x_t) p(x_t | x_{t-1}, u_{t-1}) Bel(x_{t-1})}{p(x_t | x_{t-1}, u_{t-1}) Bel(x_{t-1})} \\ &\propto p(z_t | x_t) \end{aligned}$$


Resampling

- **Given:** Set S of weighted samples.
- **Wanted** : Random sample, where the probability of drawing x_i is given by w_i .
- Typically done n times with replacement to generate new sample set S' .


Resampling




Motion Model Reminder




Proximity Sensor Model Reminder





Laser sensor





Sonar sensor



















































Sample-based Localization (sonar)



Initial Distribution




After Incorporating Ten Ultrasound Scans




After Incorporating 65 Ultrasound Scans



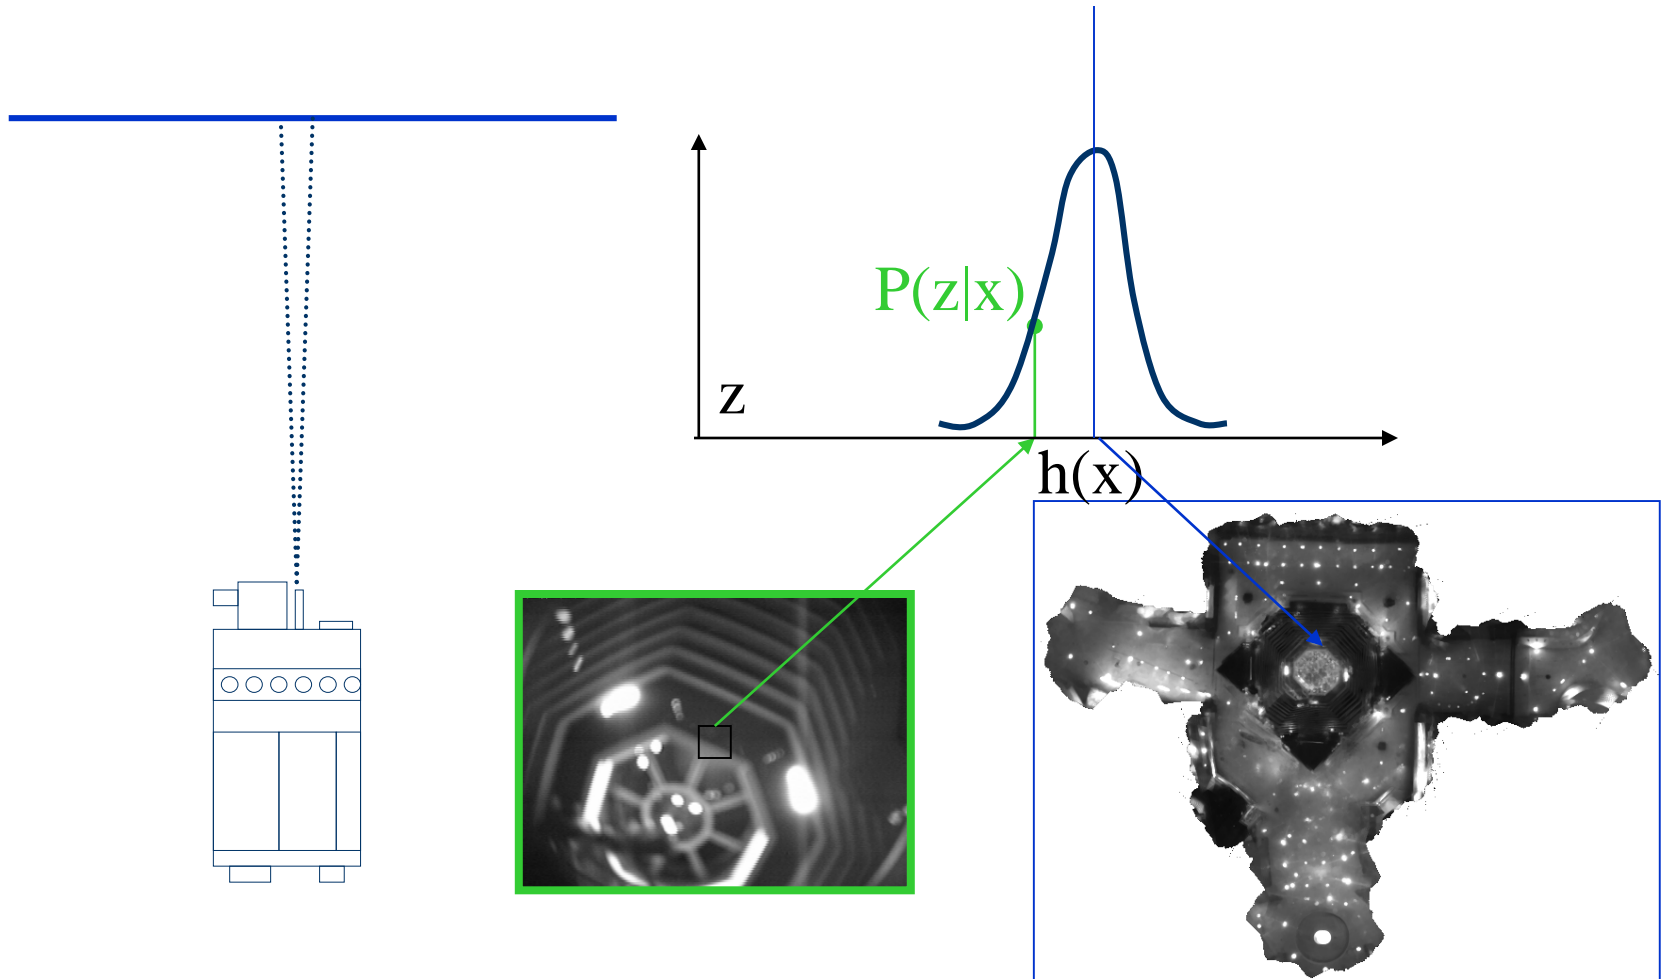
Estimated Path



Using Ceiling Maps for Localization




Vision-based Localization



Under a Light

Measurement z :




$P(z/x)$:




Next to a Light

Measurement z :



$P(z/x)$:




Elsewhere


Measurement z :



$P(z/x)$:



Global Localization Using Vision



Limitations

- The approach described so far is able to
 - track the pose of a mobile robot and to
 - globally localize the robot.
- How can we deal with localization errors (i.e., the kidnapped robot problem)?

Approaches

- Randomly insert samples (the robot can be teleported at any point in time).
- Insert random samples proportional to the average likelihood of the particles (the robot has been teleported with higher probability when the likelihood of its observations drops).

Random Samples Vision-Based Localization

936 Images, 4MB, .6secs/image

Trajectory of the robot:



Odometry Information



Resulting Trajectories

Position tracking:




Resulting Trajectories


Global localization:



Global Localization



Kidnapping the Robot



Summary

- Particle filters are an implementation of recursive Bayesian filtering
- They represent the posterior by a set of weighted samples.
- In the context of localization, the particles are propagated according to the motion model.
- They are then weighted according to the likelihood of the observations.
- In a re-sampling step, new particles are drawn with a probability proportional to the likelihood of the observation.

Reading

- “Probabilistic Robotics”, Sebastian Thrun, Wolfram Burgard, and Dieter Fox, Chapter 2, Chapter 3

Logistics

- Great job in Project 1 😊
- No lab this week 😊
- Project 2, Phase 1 due this Friday (13 Nov.)
- Project 2, Phase 2 to be released this week, due 20 Nov.