

Reinforcement Learning: Introduction

COMP4211



THE DEPARTMENT OF
COMPUTER SCIENCE & ENGINEERING
計算機科學及工程學系

Supervised learning

- the learner is provided with a set of inputs **together** with the corresponding desired outputs


Unsupervised learning

- training examples as input patterns, with **no** associated output patterns

Reinforcement learning


- Given: input and **evaluative** output only

Example (Pole balancing)




- **goal**: balance the pole as long as possible
- **states**: dynamic states of cart-pole system
- **actions**: push left, push right
- **rewards**: always 0 unless pole falls or cart hits end of track, in which case -1

Example (Mountain car)



- **goal**: minimize time to the "goal"
- **states**: car's position and velocity
- **actions**: forward, reverse, none
- **rewards**: always -1 until car reaches the goal


Example (Backgammon)



- goal: win
- states: configurations of the playing board
- actions: moves
- rewards: win: +1, lose: -1, else: 0

Example (Acrobat)


Goal: Raise tip above line



- goal: minimize time to “goal”
- state variables: 2 joint angles, 2 angular velocities
- rewards: -1 per time sweep

What is Reinforcement Learning (RL)?

Learning from interacting with an environment to achieve a goal




Learning a mapping from **states** to **actions** to maximize **long-term reward**

Given: a finite set of **states** S and a set of **actions** A

Example (Vacuum world)

Two locations, each location may or may not contain dirt, and the cleaner may be in one location or the other




- 8 possible states
- Possible actions: left, right, and suck

RL Framework...

At each discrete time, agent

- observes state $s_t \in S$ and
- chooses action $a_t \in A$

Example (Vacuum world)




RL Framework...

At each discrete time, agent

- observes state $s_t \in S$ and
- chooses action $a_t \in A$


then

- receives immediate reward r_t and
- state changes to s_{t+1}




Example

States, actions, rewards, state changes




- G: absorbing state

Markov Assumption



s_{t+1} and r_t depend **only** on **current** state and action

$$s_{t+1} = \delta(s_t, a_t) \quad \text{and} \quad r_t = r(s_t, a_t)$$




- **Markov decision process** (MDP)

Deterministic vs Non-Deterministic

Deterministic

Example



start in 1

- action “right” goes to 2

Non-deterministic: Actions may have **uncertain** outcomes

Example

action “suck” can dirty a clean carpet

- start in #4, action “suck” \rightarrow reach $\{2,4\}$

Deterministic or Non-Deterministic?

Example

chess?

Example

car driving?

Example

robotic control?


Non-Deterministic

Actions may have **uncertain** outcomes

- $P(s, s', a)$: **probability** of transition from s to s' given action a
- $R(s, s', a)$: **expected** reward on transition s to s' given action a

Example

You run a startup company. In every state you must choose between “Saving money” or “Advertising”.




Policy

Learn a mapping from states to actions


- action **policy** π : $S \rightarrow A$

Example (deterministic policy)

problem



(deterministic)
policy



Example (nondeterministic policy)




actions


	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

- terminal states: shaded squares
- reward: -1 until the terminal state is reached
- actions that would take agent off the grid leave state unchanged

Random policy




Rewards



A programmer gets paid, say 10K per month. How much, in total, will the programmer earn in his life?

- $10 + 10 + 10 + 10 + \dots = \text{infinity}$



What's wrong with this argument?

Discounted Rewards

A reward (payment) in the future is **not** worth quite as much as a reward now (e.g., because of inflation)

Example

Being promised \$10,000 next month is worth only 90% as much as receiving \$10,000 right now.

Assuming payment n months in future is worth only $(0.9)^n$ of payment now, what is the programmer's future **discounted** sum of rewards?

- (reward now) + $(0.9) \times (\text{reward in 1 time step}) + (0.9)^2 \times (\text{reward in 2 time steps}) + (0.9)^3 \times (\text{reward in 3 time steps}) + (\text{infinite sum})$


Discounted Return

γ : the **discount factor** for future rewards

$$\text{discounted return} = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

- $0 \leq \gamma < 1$
- shortsighted $0 \leftarrow \gamma \rightarrow 1$ farsighted

Example (Pole balancing)



- reward = -1 upon failure; 0 otherwise
- discounted return = $-\gamma^k$ for k steps before failure
- return is maximized by avoiding failure for as long as possible